

Bitfiltrator Extended

Isabelle Friedfeld-Gebaide

Barnard College, Columbia
University
NY, NY, USA
if2266@barnard.edu

Noam Hirschorn

Columbia University
NY, NY, USA
nyh2111@columbia.edu

Dan Ivanovich

Columbia University
NY, NY, USA
dmi2115@columbia.edu

ABSTRACT

Xilinx bitstreams are encoded in Vivado to instruct FPGAs how to run code. However, the translation from code to bitstream is not distributed by Xilinx. The Bitfiltrator paper hypothesized how to reverse engineer this process on the Xilinx UltraScale and UltraScale Plus architecture. This paper attempted to recreate what that team did, and to use the same methods to analyze the Xilinx 7-Series FPGAs, the architecture preceding UltraScale.

Analysis of the original Bitfiltrator paper and code was able to be completed in addition to most of the necessary changes to convert the methods used for Xilinx UltraScale and UltraScale Plus reverse engineering of the bitstream. The paper was able to complete an analysis of the 7-Series FPGA's bitstream and extend the reverse engineering process found in the original Bitfiltrator paper to the Artix family of 7-Series FPGAs

CCS CONCEPTS

• **Hardware** → **Programmable logic elements**; • **Security and privacy** → **Hardware reverse engineering**; **Embedded systems security**.

KEYWORDS

Bitfiltrator, Bitstream, FPGA, 7-Series, Reverse Engineering, Vivado, Xilinx

1 INTRODUCTION

An FPGA (Field Programmable Gate Array) is a programmable device which can perform hardware simulation. During compilation of a piece of code for hardware simulation of an FPGA, instructions to an FPGA for how to set each of its component parts is generated, this set of instructions is known as the bitstream. Xilinx's Vivado is currently the only way to generate a bitstream for an UltraScale/UltraScale+ FPGA [7]. Unfortunately, Xilinx does not publically release the formula by which Vivado translates user code to FPGA bitstreams.

The Bitfiltrator paper [7] analyzed various fields in order to reverse engineer the addressing format and formula of the UltraScale/UltraScale+ FPGA bitstream to match up with real components of the UltraScale/UltraScale+ FPGAs. The results indicate a potential of on-the-fly reprogramming of

FPGAs without going through the overhead of using Vivado, the tool which normally generates these bitstreams.

One benefit of using the Bitfiltrator method is that on-the-fly reprogramming of FPGAs is faster. Due to the Bitfiltrator method one can bypass fault-testing and other time consuming confirmation checks from Vivado.

An second benefit of using the Bitfiltrator method, detecting if the bitstream has been modified is easier with the Bitfiltrator method. The bitstream can be read from the FPGA as it runs, and it could then be reverse engineered to ensure that the program's functionality is still as intended, and help guarantee that no changes were made just before the bitstream was generated.

A final benefit of using the Bitfiltrator method, reverse engineering of the bitstream can allow the hiding of sensitive information. If a custom bitstream can be created, sensitive information could be placed in LUTs harder to access (whether to protect from side channels or other attacks), which could be chosen via a specially crafted bitstream.

This paper attempted to extend the Bitfiltrator methods from UltraScale/UltraScale+ FPGAs architecture to the 7-Series FPGAs architecture. The architecture of the FPGAs differs in subtle but important ways from UltraScale paper, which is what Bitfiltrator [7] worked on.

2 XILINX ARCHITECTURE OVERVIEW

The Xilinx architecture is defined in terms of a hierarchy with a series of levels. These levels each define various related sections of the very regularly organized grid structure of FPGAs. The highest level is the device, which is the overall FPGA device, the next level after is the SLR (Super Logic Region), followed by the FSR (Fabric Sub Region/Clock Region) as seen in Figure 1.

These chips are often small enough that many devices only have SLR, particularly with the 7-Series FPGAs (with a select few Virtex devices having multiple) [8]. Some of the Virtex-7 chips are sufficiently large enough that Stacked Silicon Interconnect (SSI) technology is necessary in order to prevent communication between SLRs from being too slow [8]. SSI effectively increases the overall capacity of the device by combining multiple SLRs on a passive interposer layer [8].

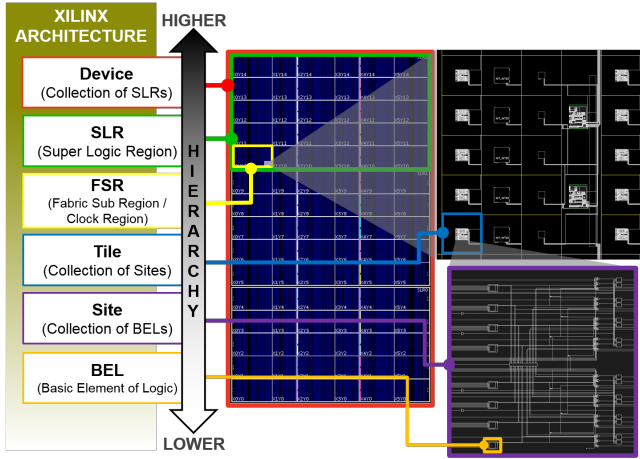


Figure 1: Levels of architectural hierarchy in Xilinx FPGAs. Starting with Device as the highest level in the architectural hierarchy in Xilinx FPGAs, followed by SLR, FSR (Clock Region), Tile, Site and ending with BEL [2].

One region relevant to the Bitfiltrator method is the Clock Region. This is a region of the SLR which uses a direct connection from the central clock of the device to serve as a clock source for all the elements in the area. It is effectively a neighborhood of an FPGA, with a defined size across architectures. One of the primary focuses of the Bitfiltrator paper and method [7] is to translate data from the clock region to its addressing location as encoded in the bitstream.

Each clock region contains rows with two types of resources – that are important to the Bitfiltrator method – CLBs (configurable logic blocks) and BRAMs (Block RAM/memory)[7]. CLBs allow for the FPGA to actually perform logic and calculations while BRAM serves as memory. These rows also have major columns and minor columns, which is the lowest level defined in the bitstream [7]. This can be compared to a building number in a mailing address, while SLRs can be compared to the state.

This lowest level of granularity defines the addresses of slices, which actually make up individual CLBs and BRAM. These resources make up the rows and columns inside of a clock region. Additionally, there is a specialized type of CLB which is a DSP (Digital Signal Processor) [5]. Each slice will help make up one of these resources.

The bitstream assigns address values to each slice. The will call a slice’s address with the FAR (Frame Address Register) and then use the FDRI (Frame Data Register Input) to pass on configuration details. The FAR is defined by 4 primary elements in the UltraScale architecture and 5 in the 7 series. The SLR element defines the SLR of the slice, while the major row, major column, and minor column provide

more detailed information. The extra element in the 7 Series is the top/bottom bit, also known as the half bit or bot. This bit determines which half of the device the slice is in, in a manner which will be in Section 2.1 [7].

The final relevant register is IDCODE, which contains the name of the SLR (which is important if there are multiple), since information is often with respect to the SLR rather than the device. For example, the major row in UltraScale is defined as the row slice in the SLR [7].

2.1 Key Differences between 7-Series FPGAs and UltraScale

Table 1: Key Addressing Differences between 7-Series and UltraScale/UltraScale+ FPGAs in Bits.

	7-Series	UltraScale	UltraScale+
Padding Size	5	5	4
Block Size	2	2	2
Bot Bit Size	1	0	0
Major Row Size	4	5	5
Major Column Size	9	9	9
Minor Column Size	6	6	7

Table 2: Key Architectural Differences between 7-Series and UltraScale/UltraScale+ FPGAs .

	7-Series	UltraScale	UltraScale+
Bot Bit [4].	Yes	No	No
Unified Clock Region	No	Yes	Yes
BRAM location [6]	18E1	18E2	18E2
DSP Register [5]	48E1	48E2	48E2

The primary difference between the 7-Series FPGAs and UltraScale/UltraScale+ architectures is the introduction of a top/bottom bit (Bot) in the FAR, as can be seen in Figure 2. In UltraScale, the rows of a clock region are simply incremented from the bottom of SLR, to create a relative row address, as can be seen in Figure 3. In the 7-Series FPGAs, the relative row address is instead calculated as a distance from the center of the device, using the top/bottom bit. The 7-Series top/bottom configuration contains an additional layer of hierarchy, and it completely changes the calculation of the relative row address when compared to its successor, the UltraScale/UltraScale+ FPGAs.

Another difference is regarding the basic 18kB BRAM block. In the UltraScale, the block is the 18E2 while the 7-Series FPGAs use the 18E1, which is instantiated differently in Verilog [6]. The 18E2 is a more updated version which




Figure 2: Xilinx 7-Series FAR (Frame Address Register) [4].




Figure 3: The frame address format in UltraScale and UltraScale+ devices. The address is composed of a 4-tuple: block type, major row, major column, and minor column.

The block type is either `CLB_IO_CLK` or `BRAM_CONTENT`. The number above/below each field marks its start/end offset in the 32-bit FAR word. UltraScale+ devices have an expanded minor column field compared to UltraScale devices [7].

has inputs to cascade between BRAM blocks, but that element was not actually used in the paper. Nevertheless, it is important to switch the block and code in `ring_bram.v` to use the type of BRAM that is actually on the device.

3 METHODOLOGY

For the most part, reversing the bitstream for a 7-Series device can be done using the same method as the original paper. The original paper’s codebase revolved around the CLI-driven `create_device_summary.py`, which uses a combination of Python, Verilog, and TCL scripts to generate and parse the necessary bitstreams.

`create_device_summary.py` exactly follows the reversing methodology described in the paper, so the team set the measurable goal of adjusting the codebase such that the steps in `create_device_summary.py` will successfully run on a 7-Series device. This process involved adjusting almost every step of the original process to account for major architectural differences — particularly the lack of SLRs, the addition of half bits, and the change in logic location formats. Ultimately, this goal was reached.

The largest changes made to adapt to the 7-Series devices was the introduction of half bits and the removal of SLRs. The data structures in the code by Kashani et al. was highly dependent on there being multiple SLRs, and major changes had to be made to virtually all functions used in the `create_device_summary.py` process to parse and store

bitstream information in the authors’ JSON-like defaultdict trees. Further changes had to be made to account for differences in metadata encoding in the bitstream, such as the FPGA part numbers. 7-Series devices store their part numbers differently, which made the Bitfiltrator code by Kashani et al. completely unable to determine the architecture of a generated 7-Series bitstream. For a basic summary of how changes regarding the top/bottom bit were implemented, see `major_architecture_level_changes.txt`.

In general, the team tried to implement all changes in a way that would still allow the codebase to function on UltraScale/UltraScale+ devices as it originally had. Sometimes, this was easy to do, but major data structure adjustments for the inclusion of half bits often made this much, much harder. For the sake of completing the adaptation in time, many of these more-difficult adjustments, which would’ve required writing completely separate classes and functions for UltraScale+ vs 7-Series, were left unimplemented. Thus, the codebase now only works for the 7-Series devices — although it will still be able to parse UltraScale+ devices correctly, the major changes made to core data structures means the modified codebase is unable to store its results in memory. Some further simple but incredibly-tedious development work would be required to make this codebase a complete extension capable of handling UltraScale, UltraScale+, and 7-Series devices alike.

In development work, Artix-7 devices were used. As the “lower power and cost” option of the 7-Series families, the

Artix-7 boards had fewer components, which made them ideal for development work. The submitted codebase has been tested and confirmed to successfully step through the `create_device_summary.py` process on multiple Artix-7 FPGAs.

3.1 Issues Faced

One difficulty of the paper was in access. There were issues in running the code on some computers, which made it much more difficult to debug issues as they primarily had to be solved through static analysis, and fixes could not be quickly and easily tested. The team therefore had times when some team members were unable to contribute even when they had time, to lack of being able to determine what the next immediate steps should be.

An additional issue was lack of clarity on the goal. Initially, there were some issues in figuring out what the project was supposed to be, whether it be to focus on verifying the findings of the original paper or attempting to recreate its methods without using the provided code. This led to a delay in actually being able to start the project proper.

The code produced by this paper does not seem to fully work on the other 7-Series families (Kintex and Vertex), as all of this team’s development time was spent working with Artix-7. There are some other small but critical issues in how some attributes of the Kintex and Vertex devices are encoded, such as DSP and LL naming, but the paper’s code is able to get devices from these families most of the way through the reversing process.

4 BREAKING BITFILTRATOR

A portion of time was spent analyzing the Bitfiltrator [7] code to see if there would be any ways to prevent this sort of reverse engineering.

Xilinx does not publicly release its methodology for creating bitstreams since it would prefer for it to remain a secret. The primary assumptions of the code seem to be on the consistency of the layout, which would be a great cost to change. Even making the size of clock regions not quite as standardized would disrupt addressing used in Bitfiltrator, but would again involve disrupting the regularity of the device.

However, an easier method Xilinx could use to stop this method would be to simply change the order/way that the code translates between a slice or clock region (as given in the absolute format of `X_Y_`) and its corresponding address in the FAR (as determined by its major row/column and minor column. Right now, Bitfiltrator [7] was able to use the information that the FAR will simply increment from the bottom row of each SLR in the UltraScale architecture.

This relationship was also explained with regards to 7-Series FPGAs, albeit following a different pattern than UltraScale. If Xilinx were to have this relationship follow some sort of encoded pattern however (so it could not be determined from the user guide what major row `Y_` would map to), this would disrupt the premise on which the Bitfiltrator method [7] was founded, with the added benefit that it would not actually require a change in hardware. Instead, all that would be necessary would be to edit Vivado so it would know how to perform this mapping. This could be viewed as a safe proposition since the point of the Bitfiltrator method [7] is that it is difficult to actually track the operations Vivado performs to transform code into a bitstream. Therefore, there is reason to believe that the encoding would be hard to break just by attempting to find it in Vivado.

An example of such a system would be a symmetric cipher like Advanced Encryption Standard (AES) [3]. The X and Y locations would be encoded and then passed in the FAR to the FPGA. The FAR would then be decoded on the FPGA to retrieve the actual major row being defined. There are of course a few advantages and disadvantages to the proposed

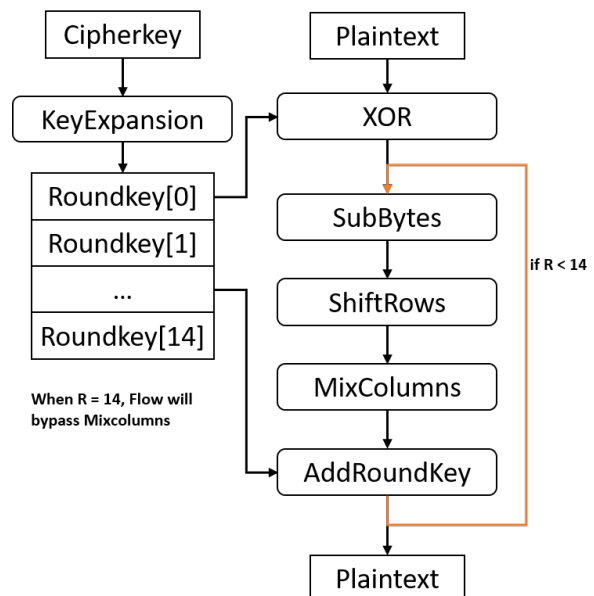


Figure 4: AES-128/192/256 encryption consist of 5 parts: KeyExpansion, SubBytes, ShiftRows, MixColumns and AddRoundKey.

AES encryption first does XOR add to input plain data blocks with first roundkey. Then AES-128/192/256 encryption performs 10/12/14 round of processing with the left round keys, each at a time. Each round sequentially does SubBytes, ShiftRows, MixColumns and AddRoundKey [1].

AES model.

The AES method would force a Bitfiltrator style attacker to break AES to pass the correct data to correspond to the major row. However, there are a few ways this approach would need to be customized to work on an FPGA for these purposes. Firstly, there are a limited number of clock regions on a device, likely not enough to fill an AES-256 cipher block. Instead, random data may need to be generated which would be marked in a way that the FPGA knew to ignore (such as being larger than 128, as there are not even close to that many clock regions per SLR). Another idea to note is that a few small numbers such as 0,1, and 2 will be repeated a lot, particularly on smaller devices with only a few clock regions. Therefore, AES may also need to be modified slightly to be keyed with some element which would change between bitstreams to prevent a brute force attack from reverse engineering the AES key by using a small device with only a few options for major row and major column values. This could actually be potentially solved by the inclusion of random numbers as suggested above.

There are a few downsides of the proposed AES method. The first is that it would force the FPGA to decode its own bitstream. The entire point of the Bitfiltrator project is that the FPGA is too slow at generating bitstreams [7] and this would add further delays. Also, a bit stream is meant to be able to easily pass information to the FPGA, while this process would make programming it more complicated. The additional hardware for this would also add expense to the actual product. Moreover, this method assumes that the FPGA decoding circuit itself could not be attacked. The FPGA has to store the AES key somewhere, and it does not have the designated memory safety storage of a normal computer.

Another simple option would be to simply restrict information about the size of different parts of the hierarchy. In particular, it is important to know the size of a clock region to then be able to measure where it is. The Bitfiltrator [7] paper gives an example of finding the SLR from knowing the number of CLBs per clock region and the number of clock regions per SLR [7]. While the number of clock regions per SLR is something the method had to figure out, a user guide explicitly states how tall each clock region is. This information is important for being able to make further inferences in the rest of the paper. There is a cost to hiding this sort of data, in that it can make it generally harder to use/study the FPGA if a user needs to use the size of the clock regions for some other reason. Nevertheless, this could also be a viable option as it avoids the extremely distasteful option of disrupting the regularity of the hardware.

Both of the methods would help prevent Bitfiltrator type attacks on Xilinx FPGAs, but in two very different ways. The first method involves using formal encryption methods to actually encode the data. This has the upside of being more

technically secure (at least from this sort of attack), but has downsides including opening up another attack surface (to get the key from the FPGA itself). Additionally, it will add overhead to programming an FPGA from the bitstream (and it should be noted that a motivation for this attack was that FPGAs already took too long to program since Vivado takes too long to generate a bitstream). Finally, the addition of this FPGA decryptor may add cost and complexity to the resulting FPGA, as it is another component to include.

The second method does not require these physical changes with all of their issues, but also is only relying on a much smaller sense of security hoping that researchers will not be able to find information which is not published (which is the sort of reverse engineering presented in the Bitfiltrator paper itself). Moreover, while the cost of restricting information may not seem like a critical burden, excluding this from the user manuals may have consequences. It could make it more difficult for users to actually analyze clock regions and SLRs if they do not know how they are organized. Additionally, less information about an FPGA generally means users are more restricted with what they can accomplish with it. Therefore, those seeking to get the most out of their device might move to a different manufacturer, hurting Xilinx's business and brand image.

5 FUTURE WORK

The changes that this paper made should be enough for the code to work on most 7-Series FPGAs. However, due to time constraints there was difficulty in finishing the debugging process. Each reference the original Bitfiltrator code made to a major row was edited to reflect the way it should be relatively addressed as outlined in the 7-Series FPGAs User guide. To test it, all that would be needed would be to go to the project U-ray comparator files and replace the U-ray database file with one from project X-ray, which is the 7-Series counterpart to project U-ray.

There is the possibility of an issue with this paper's code when dealing with a device with more than 1 SLR. In Ultra-Scale, the major rows are relative to the SLR, while in the 7-Series FPGAs, it is in reference to the distance from the center row of the device [4]. This has the potential to cause an issue if the code only finds the distance from the center of the SLR, which some parts may use instead of the center of the device due to the way the code is set up. There was a consideration to attempt to fix this error, but ultimately decided that having the code work with 1 SLR was more important, which includes most 7-Series FPGAs Products.

There is concern that more errors could occur in the more complicated task of obtaining the center of the device, due to the entire code having been built around SLRs effectively

being the highest level in the hierarchy (with all the information relevant to the SLR being entirely in the SLR array)

6 CONCLUSION

While the extension of the Bitfiltrator method to 7-Series FPGA may have run out of time, there were a few key takeaways.

First is the way different versions of an architecture can change things which do not even appear relevant at first blush. While differences such as the introduction of a top/bottom bit were expected, other issues happened simply because different versions of registers were used for the BRAM and DSP.

Another takeaway was just how similar each type of resource is in an FPGA. Everything was strictly hierarchical, and even the different resources were still treated with very similar code since they all fit into the same addressing structure with major and minor columns, major rows, half bits, and an SLR.

REFERENCES

- [1] Xilinx 2019. *AES Encryption Algorithms*. Xilinx. Retrieved May 7, 2024 from https://xilinx.github.io/Vitis_Libraries/security/2019.2/guide_L1/internals/aes.html
- [2] Xilinx 2023. *Xilinx Architecture Terminology*. Xilinx. Retrieved May 5, 2024 from https://www.rapidwright.io/docs/Xilinx_Architecture.html
- [3] Kiteworks 2024. *Everything You Need to Know About AES-256 Encryption*. Kiteworks. Retrieved May 7, 2024 from <https://www.kiteworks.com/risk-compliance-glossary/aes-256-encryption/>
- [4] Rick Altherr. 2018. *Unpacking Xilinx 7-Series Bitstreams: Part 3*. Retrieved May 5, 2024 from <https://www.kc8apf.net/2018/08/unpacking-xilinx-7-Series-bitstreams-part-3/>
- [5] AMD 2021. *UltraScale Architecture DSP Slice*. AMD. <https://docs.amd.com/v/u/en-US/ug579-ultrascale-dsp>.
- [6] AMD 2021. *UltraScale Architecture Libraries Guide (UG974)*. AMD. <https://docs.amd.com/r/2021.1-English/ug974-vivado-ultrascale-libraries/RAMB18E2>.
- [7] Sahand Kashani, Mahyar Emami, and James R. Laru. 2022. Bitfiltrator: A general approach for reverse-engineering Xilinx bitstream formats. In *2022 32nd International Conference on Field-Programmable Logic and Applications (FPL)*. 01–08. <https://doi.org/10.1109/FPL57034.2022.00039>
- [8] Xilinx 2016. *7 Series FPGAs Configurable Logic Block*. Xilinx. <https://fpga.eetrend.com/files-eetrend-xilinx/download/201408/7595-13762-ug4747seriesclb.pdf>.